SILESIAN UNIVERSITY OF TECHNOLOGY

FACULTY OF AUTOMATIC CONTROL, ELECTRONICS AND
COMPUTER SCIENCE

CONTROL, ELECTRONIC AND INFROMATION ENGINEERING

Engineering thesis

Containerization of data warehouse creation and management processes

Author: Karol Latos

Supervisor: dr inż. Anna Gorawska

Gliwice, November 2021

Contents

Chapter 1

Introduction

The subject of this engineering thesis is a standalone data warehouse built using multiple Docker containers. The thesis describes how the data is gathered, processed and analyzed, and what components make up each container, as well as their mutual coordination. Additionally, practical results of the project are presented in form of answers to real-world questions, together with steps taken to achieve performance optimization and code robustness.

The basic assumptions of the system are the following:

1. The application is portable, standalone and platform-independent.

2. The application requires the minimal amount of setup in order to run.

3. The code is robust in recovering from errors and able to run continuously.

4. Proper programming practices are followed whenever it is possible.

1.1   Scope of the work

The work done in relation to one of the containers in this project is partially based on a group project[1], concerned with the acquisition of data from a web service in order to answer simple questions with SQL queries. For that project, the author of this thesis implemented the data gathering part using Python language with Selenium module. This individual work

---

[1]Group members: Jakub Sieńko, Kacper Garcon and the author, collaborating during an university course *Data Warehouses and Data Mining Systems*

has been transfered and heavily modified in order to compose a container-based solution, with other containers written by the author from scratch for the purpose of this thesis.

In totality, the following parts were implemented:

- data gathering container,

- database manager container,

- data mining container,

- *nodejs* container with simple web application,

- logging service,

- flag management service,

- database management service,

- web scraping Selenium-based service,

- data management service.

Additionally, the author researched topics concerned with container orchestration as a simplification of system architecture; Selenium framework for dynamic web data retrieval; optimization of "big data" SQL queries and management of large datasets; web applications development in Node.js environment and other related fields.

1.2   Thesis contents

Chapter (2) is concerned with the changes in technology leading to various options of exploiting data availability on the Internet, while also deliniating the tools which are at the core of this project. It also explains the motivation for choosing such thesis topic.

Chapter (3) describes the domain of the problem with functional and non-functional requirements of the system. It also gives a high-level view of the data warehouse and the data inside it.

Chapter (4) contains technical information about the coordination of the containers, how they utilize the services and gives an overview of the data pipeline.

Chapters (5) and (6) are centered around the three main containers, taking a deep dive inside each one of them and following the data on its path from the website to market reports.

Chapter (7) reflects on the practical results generated from the data analysis, while contrasting them with initially posed questions. Additionally, the chapter describes the performance of the code under normal and abnormal conditions, and how the project changed with time.

Chapter (8) summarizes the effects achieved during the implementation of the system and formulates conclusions about possibilities of utilizing containerization to create and manage data warehouses.

Chapter 2

Project basis

"Data is the sword of the 21st century, those who wield it well, the Samurai."

— Jonatan Rosenerg, former Senior Vice President of Products, Google

In the contemporary world, the Internet facilitates the backbone of world-wide services, including e-commerce, transportation, entertainment and businesses. This relatively recent change presented an open world of possibilities, with new creative ideas emerging every year. In comparison to the reality of the past, the amount of information available at hand is unprecedently greater than ever before, constituting an advancement so significant, it can be argued to be as important as the invention of the printing press or even written language. Tapping into this immense ocean of knowledge is possible by utilizing adequate tools, like programming languages capable of producing specifically targeted scripts. This potentially is a fertile ground for automation; and to take advantage of it means to extend one's capability of processing information, possibly by orders of magnitude, to access previously unavailable knowledge.

2.1   Data on the internet

What is the history of internet? What is the state of the internet today? How good is the average bandwidth? How are machines communicating with each other and to what means? Instant availability of information is only practical for computers, need to automate the gathering and analysis of data. Enter web scraping.

2.2   Python language

Python is a high-level general-purpose programming language. It provides levels of abstraction from the machine architecture, so that the user doesn't have to worry about managing memory allocation or typing the variables. It can be used to develop applications in a myriad of domains. Its history begins in December 1989, when a Dutch programmer Guido von Rossum became working on a successor to the ABC language [3]. Released in 1991, with major consecutive versions in 2000 and 2008[1], it today became a robust language for just about anything, from statistics and modelling, to computer vision and creating web applications. Python is taught in schools as an entry-level language, while at the same time being used by NASA[2].

Python offers a variety of modules and packages (also called libraries), which are files of code written by other developers with collections of functionalities, for example to manage image files. Among the most popular libraries we have *numpy* — for managing matrices and multi-dimensional tables similar to MATLAB; *openCV* — for processing and analyzing images and videos; *requests* — a simple HTTP library for communicating with the server; and many more.

This project utilizes several crucial libraries, which are described in separate subsections below. The rest is briefly described in the following list:

- *checksumdir*: used for calculating a checksum of given directory, similar to a file checksum. It helps determine whether new data has arrived, is it complete and can it be passed further;

- *time*: used for measuring time between two points in code, sleeping (that is, stopping the execution) and getting information on the current date. Helps with scheduling the run until the day changes and measuring performance of various parts of code;

- *shutil*: used for creating and removing non-empty directories. Helps additionally isolate the shared data inside each container, so that it's invulnerable to changes in the original folder once running;

- *os*: used for managing local files and ensuring proper directory structure on the first run — creating shared folders for storing data, logs and flags.

---

[1]Python 3, current version.
[2]`https://github.com/nasa/podaacpy` is used for crucial communications with Jet Propulsion Laboratory

2.2.1   Selenium and Beautiful Soup

These libraries make it possible to connect to any website and emulate user behaviour in an automatic way. Selenium is a framework responsible for creating a webdriver (virtual, in-code browser object) and visiting requested URLs, while looking for specific elements (like buttons for expanding the page) or deciding whether the response from the page is complete or it needs more time to load. Beautiful Soup on the other hand is a simple but powerful system for navigating HTML code. It provides support for finding particular webpage elements like *div* or *span*, by requesting their *class*, *id*, or any other attribute. Every HTML page is converted into a soup object, similar to a nested dictionary, which then can be searched using provided methods.

2.2.2   Numpy, Pandas and Scikit-learn

Numpy, briefly mentioned before, is a mathematical library which is a basis for many other libraries, such as Pandas or Scikit-learn. Pandas revolves around Dataframes and Series — two- and one-dimensional data structures similar to Excel or SQL tables. One of the advantages is the speed of processing; selecting tens of thousands rows out of millions, based on some condition, is almost instantaneous. Dataframes come with set of tools for their manipulation, i.e. aggregation functions, joins, concatenation, etc. When this data is passed, scikit-learn is responsible for creating statistical models trying to find correlations, predict outcomes, as well as provide decision support based on the data. It is arguably one of the most important libraries in machine learning domain.

2.2.3   Matplotlib and Seaborn

Matplotlib and Seaborn are two libraries for data visualization, where the latter is an extension of the former. Matplotlib gives basic plotting functionalities and proves to be a robust module, which can be used on its own. However, in order to reduce the amount of code written and standardize the outcomes, Seaborn is used to provide more complex visualizations without an extensive use of Matplotlib.

2.3 Docker containers

Containerization (or OS-level virtualization) is a way of isolating resources inside an operating system without using virtual machines (VMs), to create self-enclosed, lightweight executables. The key difference from VMs is that virtualization uses a hypervisor — software that hosts guest operating systems and distributes hardware resources among them. Any process running inside a virtual machine only sees the guest operating system. Meanwhile, containerization uses only the host operating system and a container engine (e.g. *Docker*). From the point of view of a process running inside a container, the directory structure may largely differ from what user sees on the disk. Container should have only the minimal number of libraries and dependencies required to run it. Generally speaking, containerization is a paradigm for the operating system kernel to allow many isolated *user spaces* to exist in a shared environment, which translates to container sharing the filesystem with the host operating system without conflicts. However, recreating the container may mean losing all of our data and starting fresh. To avoid that, methods for data persistence are available, two of which are most popular and used in this project:

- Docker volumes — internal Docker storage, which persists removing the container image if stated explicitly. These volumes are not seen by the host OS.

- Bind mounts — specified directories inside the container, that will correspond to actual directories on the host operating system. This technique is similar to mounting an USB device in a filesystem.

The goal of containerization is to deploy applications securely and fast, without worrying about OS compatibility. The act of abstracting our software from the host operating system allows containers to be portable and stand-alone. Additionally, one can orchestrate many containers to run together and share the OS resources in order to perform a common task. This creates a perfect opportunity to use containerization for an application managing data warehouse, since the data has to go through many different stages in a pipeline, which correspond to separate processes, each inside its own container.

2.4 Data warehouses

Data warehouse is a data organization concept that originated in late 1980s in IBM [1]. Barry Devlin and Paul Murphy were trying to find a way to optimize the processing of data

from common source to different destinations, called decision support systems. These systems were composed of software taking various data as input, and producing a metric for finding solutions to a given problem. One example is using a decision support system highlighting unusual areas of a brain scan from a MRI, for faster recognition of potentially malicious changes. Before the practice of data warehousing, multiple systems had to acquire data independently from a business source, process it and then perform needed analysis. However, this approach yeilds several problems, most obvious of which is computational redundancy and consequently wasting of resources.

Devlin and Murphy's idea was to find commonalities between different decision support systems, gather all the needed data at once, process it and then store it in a ready-to-use format. This way, any application could request a specific set of data, tailored for its needs, without the need to perform heavy computation every time. These datasets are called *data marts*, and the isolation of decision support systems from their individual data sources proved to be a robust solution, that has quickly been implemented in businesses around the world.

In my project I gather the data from various pages on the card market website, then I process and save it in a *.csv* file format. These files are then read and converted to dataframes, which are converted to tables in a database. In order for another process to use the data efficiently, it is transformed into various data marts inside the database, creating a data pipeline from the source to the final program, which performs analysis. This way I'm utilizing data warehousing concepts, by collecting the data from a single source (one process responsible for data gathering) and delivering it to two destinations (web application and data miner).

2.5   Justification of the thesis topic

How do I use existing technology to solve an existing problem? What questions do I aim to answer?

Chapter 3

Specification

What is the actual example that I'm using? Describe the specification from the outside (helicopter view).

What is the result? Peek. Add the server side described.

Implementation of the thesis topic will by nature involve a containerization software (*here Docker*), which will host several subapplications written as semi-standalone scripts, each responsible for a part of the data pipeline. Using Python and its vast collection of libraries I'm handling the data scraping, intial cleaning and maintenance of the pre-stage database in compressed CSV files; as well as managing the MySQL database, creating helper tables, extracting new information and visualizing the data to answer user's questions. With JavaScript, SQL and HTML, I'm able to present the results in form of a simple web application, querying the database connected to a *Node.js* server.

3.1 Functional requirements

The system is required to do the following:

1. The system should gather cards, sellers and offers' data and keep it up to date

2. The system should allow the user to choose an expansion of cards to gather

3. The data gathering system should run continuously and gather data once a day

4. The gathering system should visit all card sites from specified expansion and save (a) card information, dynamic and static, and (b) full information about sale offers of this

card to CSV files

5. The gathering system should visit profile pages of all newly-encountered sellers and save their public data to a CSV file

6. The gathering system should keep track of the date and save the date data into a CSV file

7. The database manager system should run continuously and update the database whenever new complete batch of data has been gathered

8. The data miner system should run continuously and create data marts every hour, given the database has the newest data

9. The web application should run continuously, recovering from errors and providing the user with at least four ways of getting useful information from the data

3.2 Non-functional requirements

1. The total time of data processing should be less than 6 hours

2. The system should be standalone (bootstrapping itself from *docker-compose.yml* and code), and cross-platform compatible

3. The gathering system should adapt the requests frequency to the server condition

4. No user input is required after running the system

3.3 Problem domain

The main goal of this project is to optimize user's card buying decisions basing on the data gathered in the warehouse.

3.3.1 Trading card games

Trading card games (*TCG*), also known as collectible card games (*CCG*), are types of card games combining the elements of strategic gameplay and features of trading cards. The first TCG was released in 1993 under the name Magic: The Gathering (*MTG*). The game, released

by an eight-person company, was an overnight success, with over 10 million cards sold in just 6 weeks. Two years later, this basement business became a gaming corporation. Today, *MTG* is among the most popular TCGs with roughly 35 million players as of December 2018 [2].

During the game, the goal of each player is to reduce the opponent's life points by strategically playing cards from the hand, before the other one succeeds. However, each player can compile their own deck of 60[1] cards out of the thousands available. This makes every gameplay unique on a level which is fundamentally different from the classic cards. One doens't have to be an expert in the field to understand that the more cards one has, especially good cards, the higher the chances of winning. Thus, trading the cards becomes an aspect as crucial as the strategy itself.

One of the places where *MTG* cards can be bought is `www.cardsmarket.com/` — an online market with a myriad of cards from the game listed for sale, by users from around the world, majority of which lives in Europe. Users are divided into three categories: *Amateur*, *Professional* and *Powerseller*, depending on their setup and *modus operandi* (individuals, zealous hobbyists, card stores). To spend the least amount of money while collecting the most wanted cards, i.e. to optimize the shopping, one would have to analyze thousands of bits of data, from the average prices of all cards, to the velocity of sales and of restocking the virtual shelves.

3.3.2   Market-related entities

Each card has a name, an expansion, which is a higher-level grouping of cards, and rarity. These elements do not vary with time. However, the number of cards for sale, and automatically calculated price statistics change every day, that's why these are stored in a separate table, with new rows every day. Similarly, the seller has a name, country of residence and optionally an address, year of registration and what type of seller they are (private vs professional); while the table with sale offers stores every card offer from every seller, from each day. To add a unified time dimension to the data, a date table is created with unique *date_id* designating the day, month and year of the datapoint collection.

When the data from the card market is collected, it is stored in five CSV files and one text file. The text file, named after set expansion, contains card names in the order of first visit and it's used to maintain consistent card-to-card progression between runs. Then, each

---

[1]Some variations require a deck of 40 cards, which are selected from a random pool of cards

CSV file represents one entity: *card*, *seller* (static entities), *card_stats*, *date* and *sale_offer* (dynamic entities).

3.3.3 Problem formulation

From the point of view of a card collectioner, every potential buy should be analyzed for better offers in order to minimize losses. The main goals of the user among buying needed cards at the lowest possible price, without sacrificing the various card qualities; finding the most price-slashing, discout-oriented sellers; discovering which cards are more discounted than others, and thus make a better buy opportunity; or minimizing the shipping costs by ordering a whole set of cards from a single seller, should one exist. Therefore, the main questions posed by the user should be answered accordingly:

1. Given a non-empty list of card ids (and card qualities: language, condition, whether it is foiled), find three sellers that sell that card at the lowest price.

2. Given a non-empty list of card ids (and card qualities) rank them based on price statistics, to highlight the most discounted cards.

3. Given a non-empty list of card ids (and card qualities) select users selling the maximum amount of wanted cards, ordered by the lowest total price.

4. Given wanted card qualities, find which sellers have a high amount of discounted cards and the biggest differences between the average market prices and their prices.

3.4 Data warehouse modelling

The data warehouse lifecycle beings when the data is collected. The pre-stage area encompasses what is written in the CSV files after a gathering run. These files correspond to entities from subsection (3.3.2) and to the tables presented below.

3.4.1 Base tables

After the data is gathered from the website, and stored in the CSV files, the database manager updates an image-based MySQL server running in Docker, creating the tables corresponding to each entity (seen in Fig. 3.1)— and data miner adds new tables as data marts for future use, and performs other analyses and visualizations.
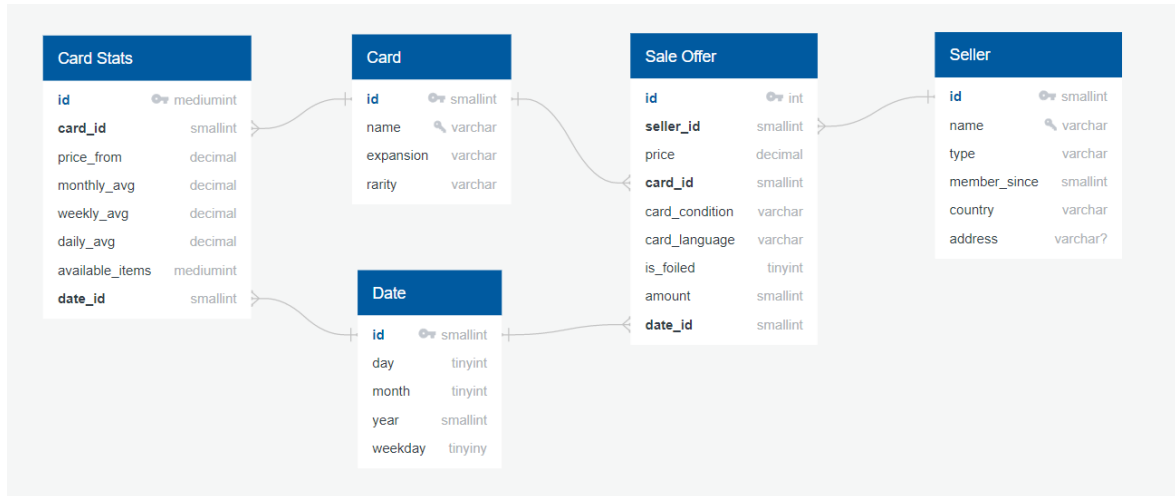
Figure 3.1: Database schema with the five base tables

3.4.2   Helper tables

The data miner creates new tables basing on the five elemental ones, in order to speed up the processing of SQL queries. The largest and most problematic table is the *sale_offer* with few millions of records, eight attributes each. To reduce the loading time, the newest day is selected as a separate table. Similarly, last two weeks are another temporal subset, extraction of which helps with performance.

3.5   Containers overview

The task of the *firefox_webdriver* container is to provide a constant and reliable access to a browser-like solution for an automated page navigation.

The *mysql_database* container hosts a database with persistent memory, storing the major part of the data warehouse and making it accessible to other containers.

The *data_gathering* subsystem automatically scrapes and collects the web data in the *data* directory, while *database_manager* takes this data and passes it to the database. It is important for the *database_manager* to be slightly delayed with respect to *data_gathering*, for reasons that are explained in section (4.1.7).

The *data_miner* and *node_server* containers are dependent only on the database, however here boot order and system state are not as important; these asynchonous containers perform their job properly when the data is ready and don't cause critial failure when it's

not, which is enough for the application to run reliably.

The *node_server* container hosts a website served by a Node.js server, that presents the user with four functional analyses of the card market.

Each container has its own virtual storage, so to perform tasks on a shared set of data, both docker volumes and data mounts are introduced. Additionally, each container exposes one port internally to a port on the host, inside a network encompassing the whole solution.

Chapter 4

Project architecture

The described solution is composed of four directories with subprojects:

- data-gathering,

- database-manager,

- data-miner,

- server,

three auto-generating directories with program-related data:

- logs,

- flags,

- data,

and a docker-compose.yml file. Each of the four directories has its own Dockerfile, and acts as a build context for that container image. Additionally, two containers are pulled from the Docker image repository, the MySQL server (*mysql/mysql-server*) and Selenium standalone Firefox webdriver (*selenium/standalone-firefox*). Figure (4.1) shows the dependencies between containers.

4.1   Containers

Issuing the `docker-compose up` command builds each local application image from a respective Dockerfile, and aranges them according to the configuration in docker-compose.yml
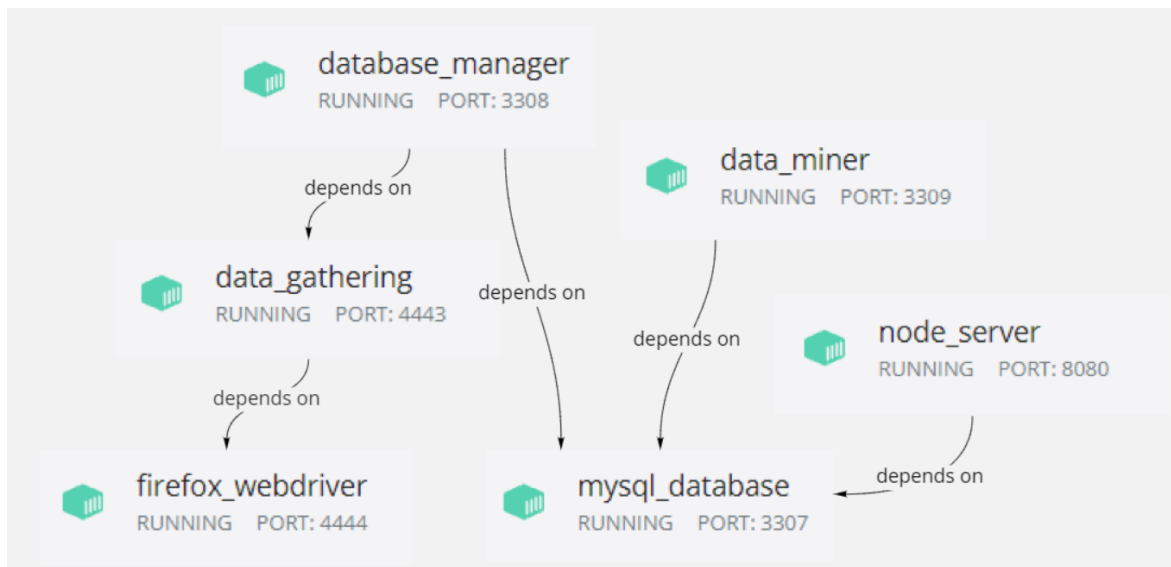
Figure 4.1: Boot order of containers (least dependant first)

file. All of the containers are booted to a shared network, with each container's name serving as an alias to be recognized immediately by other connected containers. Thus, every application exposes an internal port to the host via parameter:

```
ports:

    - "3308:3306"
```

This notation indicates that port 3308 inside the container maps to port 3306 on the host. Additionally, to persist data between system composition and decomposition, docker volumes and data binds are used:

```
volumes:

    - ./data:/data

    - ./logs:/logs

    - ./flags:/flags

    - pickle_data:/.pickles
```

Here, we declare that three directories relative to docker-compose.yml are supposed to be visible from the container's point of view as mounted in root directory. Any change to files inside the container will be immediately seen from the host level, providing the ability for several containers to work on a shared set of data. The pickles directory, on the other hand, is mapped to a *named Docker volume*, that is virtual space managed by the containerization software. These volumes are persistent and very fast, but there is no way

to peek at their contents. Only the running container can acccess the data there, therefore it acts perfectly as a space for files that should only be accessed by given container.

All of the containers also restart automatically on failure and share a set timezone. To stop the operation of the system, one can stop the app in Docker Desktop, or stop any of the containers, or delete them altogether. However, images created during the build will still exist, so to fully remove the app one can issue a command `docker-compose down -rmi local -v`, which stops and removes all containers, removes all local images, networks and volumes.

4.1.1 Container: *firefox_webdriver*

The *firefox_webdriver* container exposes port 4444 to send the traffic to the server through. It has no persistent memory and in general behaves like an abstract version of a Firefox browser running in the background, operated by any script that connects to it.

4.1.2 Container: *mysql_database*

Container *mysql_database* exposes hosts a database with persistent memory, storing the major part of the data warehouse and making it accessible to other containers.

4.1.3 Container: *data_gathering*

This subsystem automatically collects the web-scraped data in the *data* directory

4.1.4 Container: *database_manager*

This takes this data and passes it to the database. It is important for the *database_manager* to be slightly delayed with respect to *data_gathering*

4.1.5 Container: *data_miner*

These containers are dependent only on the database, however here boot order or system state are not as important; these asynchonous containers perform their job properly when the data is ready and don't cause critial failure when it's not, which is enough for the system to run reliably.

4.1.6 Container: *node_server*

These containers are dependent only on the database, however here boot order or system state are not as important; these asynchonous containers perform their job properly when the data is ready and don't cause critial failure when it's not, which is enough for the system to run reliably.

4.1.7 Container orchestration

The containers are all started when the `docker-compose up` command is issued in a root directory of the project. This causes the containers from external images to be loaded first, based on the dependency tree. From there, more control over the mutual cooridation is available via healthchecks or timing. Out of simplicity, the second approach has been selected for this thesis; i.e. the gathering project is initialized with an immediate 10 seconds delay, so that the webdriver is surely ready before a connection attempt takes place. Similarly, the database managing part is started after 20 seconds, so that the previous container can assess whether there is some new data and any action is needed, or that data can be verified, and marked as ready to be updated to the database. The data mining part is delayed 30 seconds, in case the database is in pre-update state, and other services will start to act on it. However, the most important element of orchestration is shared access to the flags directory, which stores two files with directory checksums, calculated from the ready and complete data after a gathering run. This checksum becomes the new standard and an indicator that the database contents may be outdated.

4.2 Services

The three programs written in Python share similar functions. These are grouped into **services** by their main utility and by needed external libraries. Each directory contains a *services* module with the following:

- *logs_service* — responsible for creating a container-specific logs directory with a daily log file and optionally run log file, general logging to file and console. Libraries needed: **os**, **time**.

- *flags_service* — this service creates the flags directory, sets up checksum files and provides an interface for everything related to detecting differences in directories or

checking the current stage of the data. Libraries needed: **os**, **checksumdir**.

- *data_service* — handles everything related to CSV files, translating the scraped data into attributes, loading and saving dataframes, pickling and unpickling the data, all data manipulation pre-mining and ensuring the completeness and correctness of the data. Libraries needed: **os**, **sys**, **shutil**, **time**, **pandas**.

- *web_service* — interface for connecting to the webdriver and navigating the pages, finding appropriate elements on the website, implicit and explicit waiting for response, page validation and so on. Libraries needed: **time**, **bs4**, **selenium**.

- *database_service* — manages connection to the database, setting up the tables, updating the schema with new data, as well as creating helper views and additional tables in the data mining part. Libraries needed: **mysql.connector**.

4.3   Data pipeline

High level overview of the data pipeline (schema). Where does the data come from? How do I get it? Where do I validate it? What happens in the directories? What happens in Docker volumes? What happens in the database? What is the end result of the pipeline? What is the speed of consecutive steps?

4.4   Compatilibity

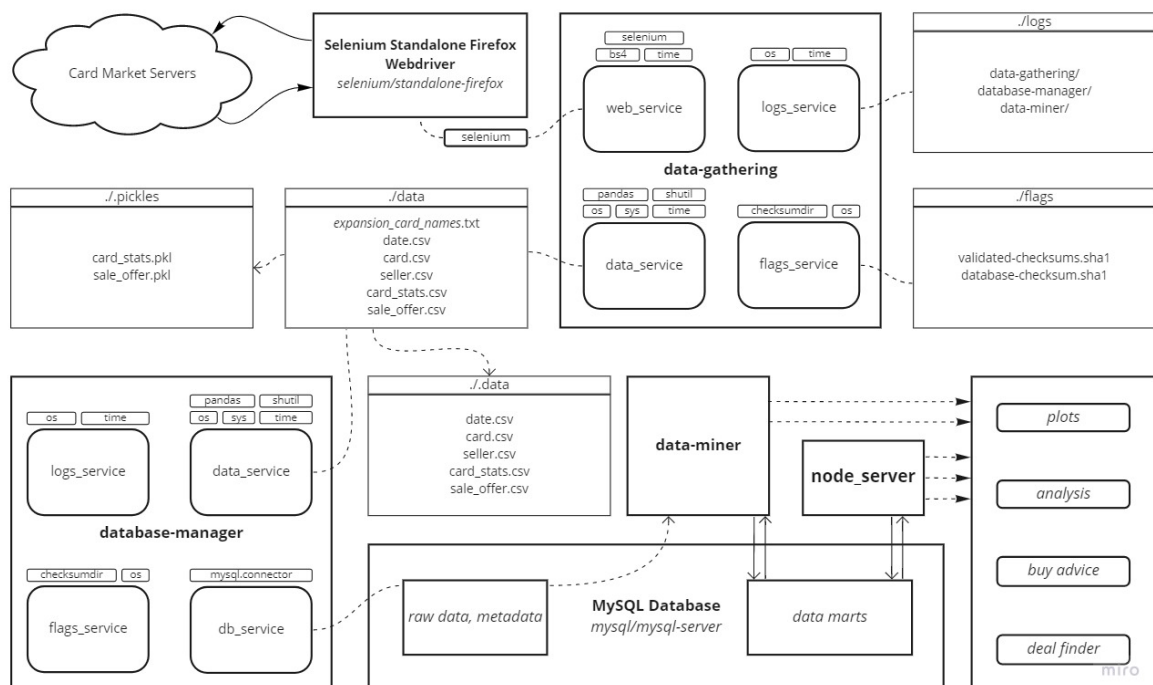Is my project compatible with main operating systems? How does the installation differ on various systems?

Figure 4.2: Data pipeline from the server (left-top) to results (right-bottom).

Chapter 5

Data gathering implementation

As mentioned before in this thesis, the *data_gathering* container consists of a Python project, which main function is to collect card market data. Static data is acquired once, dynamic data daily, and the validity of the data is ensured after every run. The user can alter the program's operation by changing values in the configuration file, and the responses are provided in console and log files. The core part of the implementation in contained in *web* and *data* services. After this container completes its objective, the CSV files in data directory are updated with the newest information available, and their checksum is calculated to be retained as a marker for the verified dataset.

5.1   Program configuration

The configurable global variables are stored in file *config.py*. There are three subgroups inside:

1. Variables for user custom program configuration.

   ```
   START_FROM = 1
   FORCE_UPDATE = False
   EXPANSION_NAME = 'Battlebond'
   ```

   The user may want to renew the daily data when suspecting incompleteness or start the operation from a specific card id (for example, starting from 128 when the previous run crashed after card 127). Moreover, the user is provided with the ability to change the expansion, so that efectively all cards from the game could be gathered.

2. Variables connected to a single run of code.

```
DATE_ID = 0
MAIN_LOGNAME = 'other_main.log'
RUN_LOGNAME = 'other_run.log'
```

These values are determined at the start of the run (whenever a new date has been detected). Log filenames are unique to the run and are used by the logging service. Shared date id is crucial for all operations of the program.

3. Fixed variables.

```
CONTAINER_DELAY = 10
NAME = 'data-gathering'
BASE_URL = 'https://www.cardmarket.com/en/Magic/Products/Singles/'
USERS_URL = 'https://www.cardmarket.com/en/Magic/Users/'
WEBDRIVER_HOSTNAME = 'firefox_webdriver'
HEADERS = {"date": {"id": "int", "day": "int", ...}, "card": ...}
```

Fixed variables are used for internal program logic. They are aggregated in a single file to be changed with convenience and to be available wherever needed by importing the file. They set up the website's paths, connection details and information about entities columns and data types.

5.2 Used services

The program uses four services:

- *Web service.* Responsible for handling the connection to the remote webdriver in another container. It serves as an interface for acquiring the cards' names from specified expansion; loading card and seller pages; clicking the *Load more* button to expand the page fully; getting website's source code in form of a Beautiful Soup object; cooling the connection down to prevent HTML response status 429[1]; and verifying the completeness of loaded pages.

- *Data service.* This biggest singular service (over 700 lines of code) utilizes the *Pandas* library to take apart soup objects and transform the information they hold into data records in csv files. This module also provides quasi-database-like functionalities,

---

[1]Too Many Requests — The user has sent too many requests in a given amount of time ("rate limiting").

like finding card's id by its name or checking whether a particular sellers or card
statistics are already saved. Here one can also find the use of *pickle* data format,
which is astonishingly fast in comparison with other data extensions, as well as the
logic responsible for pickling, unpickling, updating and verifying the data.

- *Flags service.* The main use for this service in *data_gathering* container is to signal
  whenever a dataset has been checked (collected data is tested against various incon-
  sistencies, for example the number of card statistics from each day should equal the
  number of cards, etc). When the data appears to be complete, the checksum of the
  entire data directory is calculated and saved in a shared file, for the next stages of the
  data pipeline to access and confront the datasets against.

- *Logs service.* This simple service is a wrapper for the print command. It uses the
  current time and values from the configuration file to display the state of the program
  in the console and generated log file. Each container generates a log subdirectory for
  its own purpose.

Why do I use modules as services? Why don't I use objective programming?

5.3   Local directories

The container is given access to three host directories via bind mounts: **data**, where
the csv files are stores; **logs**, with data-gathering subdirectory filled with daily and run-
level logs; and **flags**, containing SHA-1 formatted file with a list of checksums of verified
datasets.

There is also a docker volume named *pickle_data*, which maps to a folder named **.pick-
les** in the container root directory. This space is a storage for the data in pickle format,
used only by that container during a gathering run.

5.4   Run scheduling

The program begins by booting the directories up if they don't exist and determining
the current date. If it's not added to *date.csv*, new entry in the file is generated and its date
id is returned. Otherwise, just the date id is retrieved from an existing date. Then, the

execution is halted inside a scheduling loop, until the conditions for running the gathering procedure are met.

```
# Setup
logs.setup_logs()
data.setup_data()
flags.setup_flags()
data.add_date()


# Time the program execution
data.schedule_the_run()
```

As seen on Fig. 5.1, the run scheduling function checks whether the configuration flag config.FORCE_UPDATE is set or whether it's a first gathering run in the environment, which results in immediate return of control to the main function.

If that's not the case, the data directory checksum gets calculated and compared to the validated hashes. In case of occurence in the control file, the program sleeps for 60 minutes, as current data has been validated to be complete. When there is no entry in the *validated_checksums.sha1* file, the program checks whether today's data is complete. If not, the operation proceedes to gather the data, otherwise it saves the checksum as verified and waits 60 minutes. Every time after sleeping, the date is checked and date id is chosen accordingly.
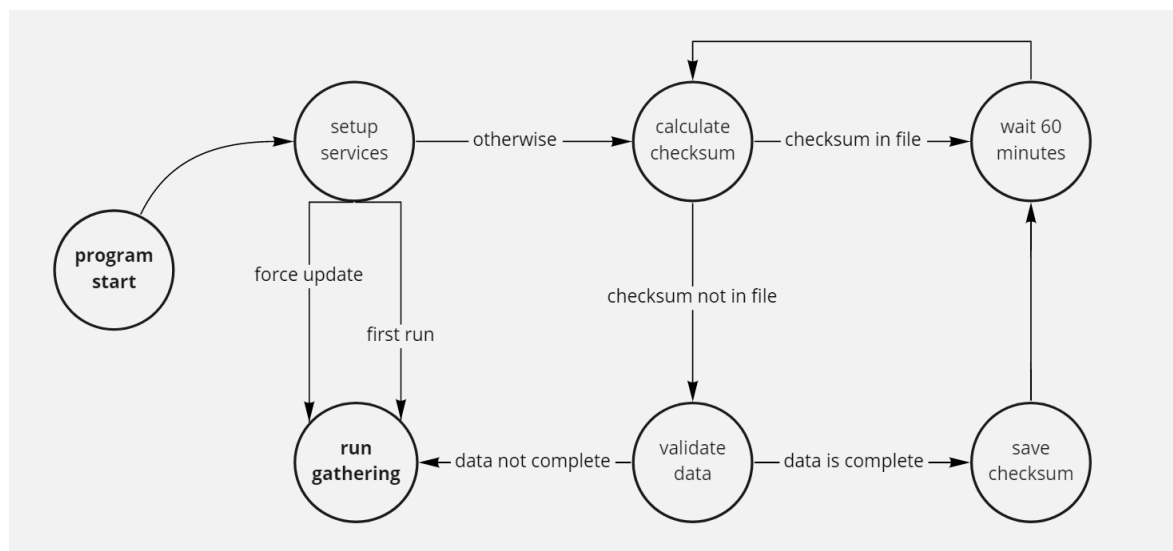

Figure 5.1: Program states in data gathering scheduling

5.5 Data pickling and validation

Initially the data was stored in raw csv files, with each new entry being appended to the end as text. This approach proved to be conceptually simple, but not very scalable, as it yielded higher and higher wait times for opening and closing files with hundreds of thousands of rows. Currently, the *sale_offer.csv* file stores roughly 10 million rows of data, but the total waiting time for opening, reading, writing and closing is greatly reduced. This is done by using the pickle data format with very fast r/w speed. Between the runs, data is kept in gzip-compressed csv format, to ensure minimal disk usage. However, as the run begins, some of the files are read into dataframes and converted to intermediate *.pkl* files. Then, whenever a pickled entity is needed, the program knows to read the pickle file instead, which loads about 100 times faster. When the program is finished, the pickled entities are transformed back to csv format. This step is the biggest hazard of potential data loss, as hours of computation kept in few, hidden files replace the previous data in a matter of minutes.

The pickle data format posed another challenge: unexpected errors concerning pickling directories and iteratively growing a pickle object. Unfortunately, the approach to acquiring sellers and cards was exactly that; a dictionary with pairs of attirbutes and values represents a single seller and to keep all entries, a list is used. Unpickling the whole object to add one seller and then pickle it back was causing the data to become unreadable after few hundret iterations. The speed of operating on gzip-compressed CSV files with less than 10 thousand rows was acceptable, especially taking into account that sellers are only added when they are not in the table, but sale offers and card stats are updated every day.

5.6 Loop over cards

The program works by querying the local TXT file with names of the cards and creating proper card urls according to a template. The webdriver loads the page with explicit wait time equal to 1.5 seconds, which gives this much time for the request elements to load. The card page features all the card information and a list of sale offers which can be expanded by clicking a "Load more" button on the bottom of the page. This button is one of the reasons for the use of Selenium framework as opposed to *requests*. We need dynamic control of the webpage to interact with the button, until all the offers are loaded. After each click, the program waits up to 1.5 seconds.
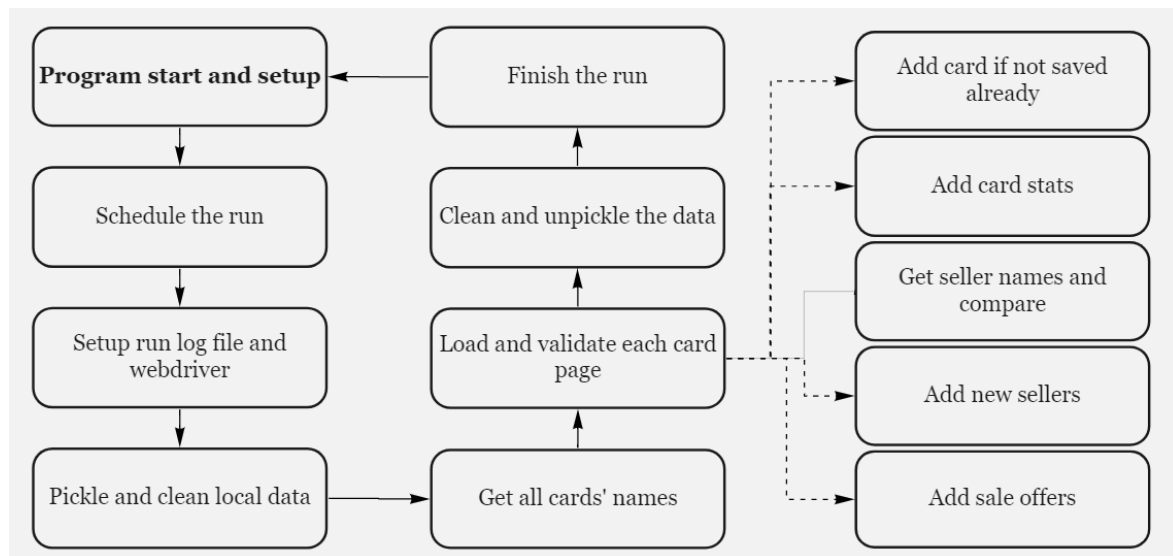
Figure 5.2: Program states in data gathering scheduling

Getting this page is the most network-heavy operation of the code, thus we save the complete webpage as a soup object to be deconstructed later. Each card soup is passed to appropriate methods to extract the card statis information and statistics, as well as the list of offers. After the card processing is completed, a START_FROM counter is incremented, so that the program can pick up from the last completed task (provided it doesn't fail into restart). After all card pages have been visited, acquired and processed, we clean and unpickle the data, then the program control comes back to the scheduler, where the data can be checked for completness and validated. (IMAGE CARD PAGE)

5.7 Soup decomposition

The card page source, packed in a soup format, is used to get the name of the card. If no such card exists in our records, it is added, together with (IMAGE LOGS ADD NEW CARD) the expansion name, rarity and a unique id. The time-related card characteristics are acquired every time. They need to be updated daily and the most recent data has priority. In case of second run of the day the saved card stats are dropped. (IMAGE CARD STATS)

5.7.1 Sale offers table

(IMAGE TABLE SALE OFFERS) Understanding the table of sale offers.

A list of sellers names if acquuired from the source code of the page, along with the list of prices, card conditions and other sale-related information. The seller names are stored in a set, from which we substract all the names of already known sellers. The remaining names are to be saved, so appropriate user page URLs are generated according to a template. Then, every new seller page is visited and data about them is saved to the seller collection. (IMAGE SELLER PAGE) Due to sellers being added cumulatively into the dataframe (use of `append(seller:  dict)`) caused pickles-related fiasco.

Updating sale offers is possible because the separate lists of seller names, prices, conditions and other, can be zipped into a dictionary. This way, we're scanning the table vertically and concatenating the attributes together to form a new set of records to be added to the sale offers file. If the offers are already saved that day, the old ones are dropped in favor of the new ones. Complete update however, happend only at the end of the run, when the pickled data is used to replace the pre-run data.

5.8   Data pipeline output

When the program is finished we have a complete list of card names, separate CSV file of static cards information with their IDs. The file with card statistics will have 264 new entries corresponding to each card and any new sellers will be saved in their CSV file. The sale offers file is going to have about 90000 more entries, representing the offers placed by sellers for each card. All the data is stored in five compressed CSV files and one TXT file. The data is ready to be checked for completion (e.g. to detect 263 cards gathered instead of 264). When the data is verified to be complete, a checksum of the data directory is calculated and saved in the validated-checksums.md5 file, so that other containers can recognize this dataset as ready to be pushed forward the data pipeline.

Chapter 6

Database management and data mining

Up to this point, our data was stored in the *data* directory and temporary pickle directory. The database managing container takes this data, ensures consistency, and uploads it to the MySQL database.

How is the data stored so far? How it can be stored? Why use database? What kinds of engines we can use? What does the data warehouse contain? What is its purpose?

6.1 Program configuration

What is in config.py file? Which values are shared among containers? Which are modifiable and which are not?

6.2 Used services

What modules are used in the subproject? What are the differences in the services?

6.2.1 Database service

What does it do, and how does it do it?

6.3 Run scheduling

How is the run scheduled? What indicators are present?

6.4 Tables from staging area

How do the tables look like when the data is gathered by the first container? What are the data types? What data can be missing or faulty? Do we have any metadata? How can it be used?

6.5 Data transformation

How will the data be transformed? What new tables will be created? How will these tables be utilized?

What information is ready to be read right away? What questions can be answered immediately? What questions are more complex? What data can be predicted from the current data?

6.6 Program configuration

What is in config.py file? Which values are shared among containers? Which are modifiable and which are not?

6.7 Used services

What modules are used in the subproject? What are the differences in the services?

6.7.1 Mining service

What does it do, and how does it do it?

6.8   Run scheduling

How is the run scheduled? What indicators are present?

6.9   Feature extraction

What new features can be discovered? `https://www.analyticsvidhya.com/blog/2021/04/guide-for-feature-extraction-techniques/`

Chapter 7

Results

What is generated by the project? What questions are answered? How does my program behave under different cases? Is it resiliant?

Check out Docker performance monitoring.

7.1 Testing

Testing of the proper functioning of each module. What is the performance of the code under normal and abnormal conditions?

7.2 Version control

How was the project developed? How were different versions maintained?

7.3 Project variation in time

How did the project evolve with time? What challenges did I encounter? How did I solve problems along the way? Which problems remained?

Chapter 8

Summary

What was the subject of this thesis? How did the assumptions or constraints influence the project? What did I implement and how did I do it?

8.1 Result assessment

What is the end result of the program? Did I manage to answer all of my questions? Are the effect satisfactory?

8.2 Scaling and customizability

How can this project be used outside of the case described in this thesis?

8.3 Conclusions

Any other closing remarks.

Bibliography

[1]   Barry A. Devlin and Paul T. Murphy. "An Architecture for a Business and Information System". In: *IBM Syst. J.* 27 (1988), pp. 60–80.

[2]   Suresh Kotha. "Wizards of the Coast". Archived from the original (PDF) on September 1, 2006, retrieved August 11, 2013. Oct. 1998. URL: `https://web.archive.org/web/20060901100217/http://faculty.bschool.washington.edu/skotha/website/cases%20pdf/Wizards%20of%20the%20coast%201.4.pdf`.

[3]   Guido Van Rossum. "The History of Python: A Brief Timeline of Python". Archived from the original on 5 June 2020. Retrieved 5 March 2021. Jan. 2009. URL: `https://python-history.blogspot.com/2009/01/brief-timeline-of-python.html`.

Appendix A

Used symbols

$M(i, j)$ — measure between points i and j.

Appendix B

Another appendix